

Getting Started in R

Aug. 2019

STAT 6000

What is R?

For now, we will use R as a fancy calculator. There are a lot of formulas in this class that would be tedious to do by hand or with a simple calculator. For instance, calculating the sample mean of a list of 100 numbers would take a while to do by hand, but seconds in R. Another convenient feature of R is its plotting capabilities. Using R one can produce clean looking plots very efficiently. Later, we will learn how to use various built in functions (such as *lm*) in R to perform data analysis.

Installation

- Installing R:
 1. Go to r-project.org.
 2. On the left side of the page, click the “CRAN” link.
 3. Scroll down to the USA section, and click a link (say, Berkeley).
 4. Click the download link for your operating system (Windows, Mac, or Linux).
 5. Follow the instructions to download the latest version of R.
- Installing RStudio (An integrated development environment for R): Go to www.rstudio.com and follow the instructions.

Using R

Assignment and Arithmetic

Assignment of variables means that in R you can create your own variables, vectors, matrices, etc. R can perform basic arithmetic operations with numbers or declared variables. Here are some examples:

```
#define variables x and y
> x = 4
> y = 3
```

```
#perform arithmetic on variables as well as numbers
```

```
> x * y - 4
```

```
#output
```

```
[1] 8
```

R can perform arithmetic on numbers, but it is really designed to work with *vectors*. Vectors are just lists of numbers, and in R we can define and perform operations on them. To create a vector, we use the concatenate function `c()`:

```
#use concatenate function to make a vector called 'myvec'
```

```
> myvec = c(4,5,6)
```

```
#display myvec
```

```
> myvec
```

```
[1] 4 5 6
```

Once we have a vector, we can perform arithmetic on the entire vector at once. We can also add two vectors together.

```
#perform arithmetic on entire vectors at once
```

```
> myvec + 3
```

```
[1] 7 8 9
```

```
> myvec * 2
```

```
[1] 8 10 12
```

```
#create another vector called 'myvec2'
```

```
> myvec2 = c(1,2,1)
```

```
#add my vectors together
```

```
> myvec + myvec2
```

```
[1] 5 7 7
```

There are built in functions that we can perform on a vector as well.

```
#sum the elements of a vector
```

```
> sum(myvec)
```

```
[1] 15
```

```
#take the product of all elements of a vector
```

```
> prod(myvec)
```

```
[1] 120
```

```
#calculate the sample mean of a vector
```

```
> mean(myvec)
```

```
[1] 5
```

```
#calculate the median of a vector
```

```
> median(myvec)
```

```
[1] 5
```

```
#calculate the sample standard deviation of a vector
```

```
> sd(myvec)
```

```
[1] 1
```

As an example, consider a simple regression case.

```
#define vectors X and Y
```

```
> X = c(1.86,.22,3.55,3.29,1.25)
```

```
> Y = c(3.34,1.79,5.66,5.83,4.74)
```

```
#Xbar is the sample mean of X
```

```
> Xbar = mean(X)
```

```
> Xbar
```

```
[1] 2.034
```

```
#similarly for Y
```

```
> Ybar = mean(Y)
```

```
> Ybar
```

```
[1] 4.272
```

```
#more advanced calculations (^ means exponent)
```

```
> sum((X - Xbar)^2)
```

```
[1] 7.81132
```

```
> sum((X - Xbar)*(Y - Ybar))
```

```
[1] 8.35866
```

Here we have everything calculated, so we can easily obtain

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} = \frac{8.35866}{7.81132} = 1.07,$$
$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X} = 4.272 - 1.07(2.034) = 2.09.$$

```
> beta1 = 8.35866 / 7.81132
> beta0 = Ybar - beta1*Xbar
```

We can now use this to find the fitted values \hat{Y}_i and residuals e_i :

```
### fitted values
> Yhat = beta0 + beta1*X
> Yhat
[1] 4.085808 2.330893 5.894226 5.616008 3.433065

## residuals
> e=Y-Yhat
> e
[1] -0.7458078 -0.5408928 -0.2342263  0.2139920  1.3069350

### check properties of residuals:
> sum(e)
[1] 3.108624e-15
> sum(X*e)
[1] 5.925815e-15

> sum(Yhat*e)
[1] 1.24345e-14
```


Plotting

Now that we have our data, and the fitted regression line, we would like to plot them both to see what we have done. The plot function in R is fairly simple. There are several options one can use when plotting, but we start with the basic ones:

```
#main is the title of the plot, col is the color of the points
> plot(X,Y,main='My Plot',col='blue')

#use abline to meake a line
> abline(beta0,beta1)
```

Here, the function `abline()` creates a line with a given slope and intercept. Here is the output for this simple example:



Example: Inference

Suppose a sequence of observations x_1, \dots, x_{50} are drawn from a normal distribution $N(\mu, \sigma^2)$. We'd like to create a 90% confidence interval for the mean parameter μ . We use the formula given in our preliminary statistics courses:

$$\bar{x} \pm z_{(1-\alpha/2)} \frac{s}{\sqrt{n}}.$$

Since $\alpha = 1 - .9 = .1$ and $n = 50$, here

$$\bar{x} \pm z_{(.95)} \frac{s}{\sqrt{50}}.$$

What we need to do now is to calculate \bar{x} , s and the appropriate z-quantile. This can be done in R:

```
#define vector x

#Compute xbar: the sample mean of x
xbar = mean(x)
xbar
[1] 0.6345216

#Compute s: the sample standard deviation of x
s = sd(x)
s
```

```
[1] 2.131412
```

```
#Compute the confidence interval
```

```
#lower bound
```

```
xbar - qnorm(.95)*s/sqrt(50)
```

```
[1] 0.138718
```

```
#upper bound
```

```
xbar + qnorm(.95)*s/sqrt(50)
```

```
[1] 1.130325
```

This gives our 90% confidence interval: (.139,1.130).

Interpretation of 90% confidence interval: If repeated samples are taken from the same population and the 90% confidence interval are computed for each sample, 90% of the intervals would contain the true population parameter.

Note that the function `qnorm()` gives the percentiles (a.k.a. quantiles) of the standard normal distribution if you provide the percentile you want. Sometimes we would like to find the percentiles for a t distribution or a Chisquare distribution or an F distribution, these are by the functions `qt()`, `qchisq()`, `qf()`, respectively.

```
##95% percentile of Chisquare distribution with 3 degrees of freedom
```

```
> qchisq(0.95, 3)
```

```
[1] 7.814728
```

```
### 95% percentile of t-distribution with 3 degrees of freedom
```

```
> qt(0.95, 3)
```

```
[1] 2.353363
```

```
## 95% percentile of F-distribution with (1,3) degrees of freedom
```

```
> qf(0.95, 1,3)
```

```
[1] 10.12796
```

We would also like to test whether the mean parameter μ is 0. We'd like to test

$$H_0 : \mu = 0 \text{ vs } H_1 : \mu \neq 0.$$

Z-statistic:

$$Z = \frac{\bar{x} - \mu_0}{s/\sqrt{n}} = \frac{\bar{x}}{s/\sqrt{n}}.$$

Given the significance level α , we reject H_0 if $|Z| > z_{(1-\alpha/2)}$, or equivalently, if $\text{pvalue} := P(|z| > |Z|) < \alpha$ (z is a standard normal random variable).

Take $\alpha = .05$, here, we reject H_0 if $|Z| > z_{(.975)}$, or equivalently, if $\text{pvalue} := P(|z| > |Z|) < .05$.

In R, we can do the hypothesis testing efficiently:

```
#Compute Z statistic
Z = xbar/(s/sqrt(50))
Z
[1] 2.105058
```

```
#Compute the critical value: 97.5% percentile of standard normal distribution
c_val = qnorm(.975)
c_val
[1] 1.959964
```

```
#Compute the p-value
p_val = 2*(1-pnorm(Z))
p_val
[1] 0.0352863
```

Exercises: Calculate the 0.975 percentile for $N(0, 1)$ – standard normal, $t_{(10)}$ – t-distribution with 10 degrees of freedom, $\chi^2_{(10)}$ – Chisquare distribution with 10 degrees of freedom, $F_{1,10}$ – F-distribution with (1, 10) degrees of freedom.